



// 15-MINUTE TALK

The Dark Sides of AI

Six ways AI quietly goes wrong — and how to keep yourself in charge.

Charis Dang and our beloved Claude · Hanoi, July 2026

// BEFORE WE START

WARM-UP_

Four quick questions. No wrong answers — hands up, shout out.

QUESTION 01 / 04

How often do you use **AI**?

A Every day
it's basically a colleague

B A few times a week

C Once in a while

D "What's ChatGPT?"

> show of hands ■

QUESTION 02 / 04

Q2

**Name some AI models —
which one is your **favorite**?**

GPT-3.5 Claude Sonnet Gemini

Claude Fable Haiku DeepSeek Claude Opus

Qwen Gemma Llama Grok Mistral GPT-4o ...yours?

```
> shout them out first - then reveal
```

Q3

QUESTION 03 / 04

How much would you **spend** on AI per month?

A

\$0

free tier forever

B

Up to \$20

one subscription

C

\$20 – \$100

power user

D

Whatever my company pays

the smartest answer

> show of hands

QUESTION 04 / 04

What's your **AI level**?

Beginner

Chatbot — ask, copy, done.



Intermediate

Claude Code, GPT Pro, skills.



Pro

MCP, multi-agent workflows.



AI Expert

Local LLMs, build your own RAG...



> whatever your level – today's stories apply to everyone ■

A brilliant new colleague...

- ◆ **Fast.** Drafts in seconds what used to take hours.
- ◆ **Confident.** Always has an answer, never hesitates.
- ◆ **Tireless.** Available 24/7, never complains.
- ◆ **...and occasionally, confidently wrong.** And it sounds *exactly the same* either way.

→ The gap between how much we use AI and how much we check it — that's where every story today lives.

// AGENDA

Six dark sides

01  Hallucination

02  Bias & unfair outcomes

03  Knowledge cut-off

04  Context loss

05  Manipulation

06  Loss of accountability

Hallucination — plausible ≠ true

A lawyer filed a brief with **six case citations from ChatGPT** — names, docket numbers, quotes. All six cases **did not exist**. He even asked ChatGPT “are these real?” — it said yes.

[Mata v. Avianca, New York, 2023 – lawyers sanctioned and fined]

Air Canada's chatbot **invented a refund policy**. A tribunal made the airline honor it.

[Moffatt v. Air Canada, 2024]

→ **AI predicts the most plausible next word — it doesn't know facts. It never says “I don't know” unless you make it.**



Zero-trust the **output**



Ask for sources

Then check they actually exist —
and say what the AI claims.



Ground it

“Answer **only** from this
document.”



Match the stakes

Brainstorm → glance. Legal,
medical, money → expert review.

→ **Confidence is not accuracy.**

Bias — prejudice at **machine scale**

AI learned from **our internet** — stereotypes included.

Amazon's hiring AI penalized the word "**women's**" on resumes. They couldn't fix it — they scrapped it.

[Reuters, 2018]

→ **AI mass-produces our bias.**



// DARK SIDE 2 • GUARDRAILS


Test it like you **don't trust it**

A/B your prompts: swap one detail. The answer shouldn't change.

Humans decide about humans

Unfair output? **Report it**

→ Same question, different name → same answer.



// DARK SIDE 3

Knowledge cut-off — frozen in time

It **fell asleep** months (sometimes years) ago — and never mentions the nap. The gap between “trained” and “today” has never closed:



→ Ask it: “what's your cutoff date?” — then check live sources for anything time-sensitive.

Context loss — goldfish memory

Message 1: “under 200 words, no competitor names.”

Message 40: it forgot both.

Context window = its short-term memory

~128K–1M **tokens** — then it overflows

No long-term memory by default

→ Long chat ≠ long memory. Restart fresh, re-state the rules, recap — don't argue.



Manipulation — nudged **against yourself**


Is this optimized **for your benefit** — or for **someone else's metric**?

Outrage keeps you scrolling

Offers timed to **weak moments**

Friction where you cancel

→ **Consent, opt-out, audit** — never target the vulnerable.



who's pulling?

// DARK SIDE 6

“The AI said so” is **not** a defense

Air Canada: “**the chatbot is its own legal entity.**”

The tribunal: no.

[Moffatt v. Air Canada, 2024]

A **named human** owns it

Can't explain it? **Not reviewed**

Spot-check, don't rubber-stamp

→ **Delegate the work — never the responsibility.**

"blame the bot?"

**YOU DEPLOY IT,
YOU OWN IT.**

// ONE MORE THING • THE HUMAN TOLL

It affects **your brain** too



Brain Fry

AI does the work — you do the **non-stop verifying**. Sustained vigilance is exhausting. MIT (2025) calls it “**cognitive debt**.”



Illusion of Expertise

Fluent answers **feel** like understanding. You can present it — but can you survive **one follow-up question**?

→ **The test: explain it out loud, without the AI. AI can do the work — it can't do the knowing.**

// THE COMMON THREAD

**Every failure = a human who
trusted without verifying.**

Every output is a **draft**

Verify by **stakes**

Your name on your work

AI drafts. Humans decide.

Thank **you**

Questions?

THE DARK SIDES OF AI · AI DRAFTS, HUMANS DECIDE